
Algorithm 1: Randomized Ensembled Double Q-Learning (REDQ)

Input: Initial policy parameters θ , ensemble of M Q-functions $\{\phi_i\}_{i=1}^M$, target Q-function parameters $\{\bar{\phi}_i \leftarrow \phi_i\}_{i=1}^M$, polyak coefficient τ , number of critics sampled $N < M$, batch size B , number of critic updates per iteration G , replay buffer \mathcal{D}

```
1 for each iteration do
2   Reset environment and observe initial state  $s_0$ ;
3   while not terminal do
4     Select action  $a_t \sim \pi_\theta(\cdot|s_t)$ ;
5     Execute  $a_t$  in environment;
6     Observe reward  $r_t$ , next state  $s_{t+1}$ , done signal  $d_t$ ;
7     Store  $(s_t, a_t, r_t, s_{t+1}, d_t)$  in  $\mathcal{D}$ ;
8     if  $d_t$  is True then
9       | Reset environment and observe new initial state  $s_{t+1}$ ;
10    end
11     $s_t \leftarrow s_{t+1}$ ;
12  end
13  for  $g = 1$  to  $G$  do
14    Sample a batch of  $B$  transitions  $(s, a, r, s', d)$  from  $\mathcal{D}$ ;
15    Sample a random subset  $\mathcal{I} \subset \{1, \dots, M\}$  of size  $N$ ;
16    Compute target Q-values:
      
$$y = r + \gamma(1-d) \left( \min_{i \in \mathcal{I}} Q_{\bar{\phi}_i}(s', a') - \alpha \log \pi_{\bar{\theta}}(a'|s') \right), \quad a' \sim \pi_{\bar{\theta}}(\cdot|s')$$

17    for  $i = 1$  to  $M$  do
18      | Update critic  $i$  parameters by gradient descent:
      
$$\phi_i \leftarrow \phi_i - \lambda_Q \nabla_{\phi_i} \frac{1}{B} \sum (Q_{\phi_i}(s, a) - y)^2$$

19      | Update target critic network  $i$ :
      
$$\bar{\phi}_i \leftarrow \tau \phi_i + (1 - \tau) \bar{\phi}_i$$

20    end
21  end
22  Update actor parameters by gradient ascent:
      
$$\theta \leftarrow \theta + \lambda_\pi \nabla_\theta \frac{1}{B} \sum \left( \frac{1}{M} \sum_{i=1}^M Q_{\phi_i}(s, \pi_\theta(s)) - \alpha \log \pi_\theta(\pi_\theta(s)|s) \right)$$

23 end
```
